

```
import warnings
warnings.filterwarnings("ignore")

import re
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.datasets import load_breast_cancer
from medmnist import BreastMNIST, INFO
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.feature_selection import SelectKBest, f_classif,
mutual_info_classif
from sklearn.impute import SimpleImputer
from sklearn.pipeline import Pipeline

from scipy import stats
from scipy.datasets import electrocardiogram
from scipy.signal import butter, filtfilt

from scipy.ndimage import (
    median_filter,
    gaussian_filter,
    sobel,
    rotate,
    shift,
    zoom
)

from sklearn.feature_selection import (
    SelectKBest,
    f_classif,
    mutual_info_classif
)

from datasets import load_dataset

from collections import Counter

from PIL import Image

sns.set_theme(style="whitegrid")

RANDOM_STATE = 42
np.random.seed(RANDOM_STATE)
```

Tabular Data

```
breast_cancer = load_breast_cancer(as_frame=True)

tabular_df = breast_cancer.frame.copy()

# Rename target to something easier to understand
tabular_df = tabular_df.rename(columns={"target": "diagnosis"})

# sklearn encoding:
# 0 = malignant
# 1 = benign
tabular_df["diagnosis_label"] = tabular_df["diagnosis"].map({
    0: "malignant",
    1: "benign"
})

print("Dataset shape:", tabular_df.shape)

tabular_df.head()
```

Dataset shape: (569, 32)

	mean radius	mean texture	mean perimeter	mean area	mean
smoothness \					
0	17.99	10.38	122.80	1001.0	
0.11840					
1	20.57	17.77	132.90	1326.0	
0.08474					
2	19.69	21.25	130.00	1203.0	
0.10960					
3	11.42	20.38	77.58	386.1	
0.14250					
4	20.29	14.34	135.10	1297.0	
0.10030					

	mean compactness	mean concavity	mean concave points	mean
symmetry \				
0	0.27760	0.3001	0.14710	
0.2419				
1	0.07864	0.0869	0.07017	
0.1812				
2	0.15990	0.1974	0.12790	
0.2069				
3	0.28390	0.2414	0.10520	
0.2597				
4	0.13280	0.1980	0.10430	
0.1809				

mean fractal dimension	...	worst perimeter	worst area	worst
------------------------	-----	-----------------	------------	-------

smoothness \				
0	0.07871	...	184.60	2019.0
0.1622				
1	0.05667	...	158.80	1956.0
0.1238				
2	0.05999	...	152.50	1709.0
0.1444				
3	0.09744	...	98.87	567.7
0.2098				
4	0.05883	...	152.20	1575.0
0.1374				

	worst compactness	worst concavity	worst concave points	worst
symmetry \				
0	0.6656	0.7119	0.2654	
0.4601				
1	0.1866	0.2416	0.1860	
0.2750				
2	0.4245	0.4504	0.2430	
0.3613				
3	0.8663	0.6869	0.2575	
0.6638				
4	0.2050	0.4000	0.1625	
0.2364				

	worst fractal dimension	diagnosis	diagnosis_label
0	0.11890	0	malignant
1	0.08902	0	malignant
2	0.08758	0	malignant
3	0.17300	0	malignant
4	0.07678	0	malignant

[5 rows x 32 columns]

```
print("FIRST FIVE ROWS")
display(tabular_df.head())

print("\nDATA TYPES")
display(tabular_df.dtypes)

print("\nDATASET INFORMATION")
tabular_df.info()

print("\nSUMMARY STATISTICS")
display(tabular_df.describe().T)
```

FIRST FIVE ROWS

	mean radius	mean texture	mean perimeter	mean area	mean
smoothness \					
0	17.99	10.38	122.80	1001.0	

0.11840				
1	20.57	17.77	132.90	1326.0
0.08474				
2	19.69	21.25	130.00	1203.0
0.10960				
3	11.42	20.38	77.58	386.1
0.14250				
4	20.29	14.34	135.10	1297.0
0.10030				
mean compactness mean concavity mean concave points mean				
symmetry \				
0	0.27760	0.3001		0.14710
0.2419				
1	0.07864	0.0869		0.07017
0.1812				
2	0.15990	0.1974		0.12790
0.2069				
3	0.28390	0.2414		0.10520
0.2597				
4	0.13280	0.1980		0.10430
0.1809				
mean fractal dimension ... worst perimeter worst area worst				
smoothness \				
0	0.07871	...	184.60	2019.0
0.1622				
1	0.05667	...	158.80	1956.0
0.1238				
2	0.05999	...	152.50	1709.0
0.1444				
3	0.09744	...	98.87	567.7
0.2098				
4	0.05883	...	152.20	1575.0
0.1374				
worst compactness worst concavity worst concave points worst				
symmetry \				
0	0.6656	0.7119		0.2654
0.4601				
1	0.1866	0.2416		0.1860
0.2750				
2	0.4245	0.4504		0.2430
0.3613				
3	0.8663	0.6869		0.2575
0.6638				
4	0.2050	0.4000		0.1625
0.2364				
worst fractal dimension diagnosis diagnosis_label				

0	0.11890	0	malignant
1	0.08902	0	malignant
2	0.08758	0	malignant
3	0.17300	0	malignant
4	0.07678	0	malignant

[5 rows x 32 columns]

DATA TYPES

mean radius	float64
mean texture	float64
mean perimeter	float64
mean area	float64
mean smoothness	float64
mean compactness	float64
mean concavity	float64
mean concave points	float64
mean symmetry	float64
mean fractal dimension	float64
radius error	float64
texture error	float64
perimeter error	float64
area error	float64
smoothness error	float64
compactness error	float64
concavity error	float64
concave points error	float64
symmetry error	float64
fractal dimension error	float64
worst radius	float64
worst texture	float64
worst perimeter	float64
worst area	float64
worst smoothness	float64
worst compactness	float64
worst concavity	float64
worst concave points	float64
worst symmetry	float64
worst fractal dimension	float64
diagnosis	int64
diagnosis_label	str
dtype: object	

DATASET INFORMATION

<class 'pandas.DataFrame'>

RangeIndex: 569 entries, 0 to 568

Data columns (total 32 columns):

#	Column	Non-Null Count	Dtype
0	mean radius	569 non-null	float64
1	mean texture	569 non-null	float64
2	mean perimeter	569 non-null	float64
3	mean area	569 non-null	float64
4	mean smoothness	569 non-null	float64
5	mean compactness	569 non-null	float64
6	mean concavity	569 non-null	float64
7	mean concave points	569 non-null	float64
8	mean symmetry	569 non-null	float64
9	mean fractal dimension	569 non-null	float64
10	radius error	569 non-null	float64
11	texture error	569 non-null	float64
12	perimeter error	569 non-null	float64
13	area error	569 non-null	float64
14	smoothness error	569 non-null	float64
15	compactness error	569 non-null	float64
16	concavity error	569 non-null	float64
17	concave points error	569 non-null	float64
18	symmetry error	569 non-null	float64
19	fractal dimension error	569 non-null	float64
20	worst radius	569 non-null	float64
21	worst texture	569 non-null	float64
22	worst perimeter	569 non-null	float64
23	worst area	569 non-null	float64
24	worst smoothness	569 non-null	float64
25	worst compactness	569 non-null	float64
26	worst concavity	569 non-null	float64
27	worst concave points	569 non-null	float64
28	worst symmetry	569 non-null	float64
29	worst fractal dimension	569 non-null	float64
30	diagnosis	569 non-null	int64
31	diagnosis_label	569 non-null	str

dtypes: float64(30), int64(1), str(1)

memory usage: 146.4 KB

SUMMARY STATISTICS

	count	mean	std	min	\
mean radius	569.0	14.127292	3.524049	6.981000	
mean texture	569.0	19.289649	4.301036	9.710000	
mean perimeter	569.0	91.969033	24.298981	43.790000	
mean area	569.0	654.889104	351.914129	143.500000	
mean smoothness	569.0	0.096360	0.014064	0.052630	
mean compactness	569.0	0.104341	0.052813	0.019380	
mean concavity	569.0	0.088799	0.079720	0.000000	
mean concave points	569.0	0.048919	0.038803	0.000000	
mean symmetry	569.0	0.181162	0.027414	0.106000	
mean fractal dimension	569.0	0.062798	0.007060	0.049960	

radius error	569.0	0.405172	0.277313	0.111500
texture error	569.0	1.216853	0.551648	0.360200
perimeter error	569.0	2.866059	2.021855	0.757000
area error	569.0	40.337079	45.491006	6.802000
smoothness error	569.0	0.007041	0.003003	0.001713
compactness error	569.0	0.025478	0.017908	0.002252
concavity error	569.0	0.031894	0.030186	0.000000
concave points error	569.0	0.011796	0.006170	0.000000
symmetry error	569.0	0.020542	0.008266	0.007882
fractal dimension error	569.0	0.003795	0.002646	0.000895
worst radius	569.0	16.269190	4.833242	7.930000
worst texture	569.0	25.677223	6.146258	12.020000
worst perimeter	569.0	107.261213	33.602542	50.410000
worst area	569.0	880.583128	569.356993	185.200000
worst smoothness	569.0	0.132369	0.022832	0.071170
worst compactness	569.0	0.254265	0.157336	0.027290
worst concavity	569.0	0.272188	0.208624	0.000000
worst concave points	569.0	0.114606	0.065732	0.000000
worst symmetry	569.0	0.290076	0.061867	0.156500
worst fractal dimension	569.0	0.083946	0.018061	0.055040
diagnosis	569.0	0.627417	0.483918	0.000000
		25%	50%	75%
max				
mean radius	11.700000	13.370000	15.780000	
28.11000				
mean texture	16.170000	18.840000	21.800000	
39.28000				
mean perimeter	75.170000	86.240000	104.100000	
188.50000				
mean area	420.300000	551.100000	782.700000	
2501.00000				
mean smoothness	0.086370	0.095870	0.105300	
0.16340				
mean compactness	0.064920	0.092630	0.130400	
0.34540				
mean concavity	0.029560	0.061540	0.130700	
0.42680				
mean concave points	0.020310	0.033500	0.074000	
0.20120				
mean symmetry	0.161900	0.179200	0.195700	
0.30400				
mean fractal dimension	0.057700	0.061540	0.066120	
0.09744				
radius error	0.232400	0.324200	0.478900	
2.87300				
texture error	0.833900	1.108000	1.474000	
4.88500				
perimeter error	1.606000	2.287000	3.357000	

21.98000			
area error	17.850000	24.530000	45.190000
542.20000			
smoothness error	0.005169	0.006380	0.008146
0.03113			
compactness error	0.013080	0.020450	0.032450
0.13540			
concavity error	0.015090	0.025890	0.042050
0.39600			
concave points error	0.007638	0.010930	0.014710
0.05279			
symmetry error	0.015160	0.018730	0.023480
0.07895			
fractal dimension error	0.002248	0.003187	0.004558
0.02984			
worst radius	13.010000	14.970000	18.790000
36.04000			
worst texture	21.080000	25.410000	29.720000
49.54000			
worst perimeter	84.110000	97.660000	125.400000
251.20000			
worst area	515.300000	686.500000	1084.000000
4254.00000			
worst smoothness	0.116600	0.131300	0.146000
0.22260			
worst compactness	0.147200	0.211900	0.339100
1.05800			
worst concavity	0.114500	0.226700	0.382900
1.25200			
worst concave points	0.064930	0.099930	0.161400
0.29100			
worst symmetry	0.250400	0.282200	0.317900
0.66380			
worst fractal dimension	0.071460	0.080040	0.092080
0.20750			
diagnosis	0.000000	1.000000	1.000000
1.00000			

```

missing_values = pd.DataFrame({
    "Missing Values": tabular_df.isnull().sum(),
    "Percentage": (tabular_df.isnull().sum() / len(tabular_df)) * 100
})

display(
    missing_values[
        missing_values["Missing Values"] > 0
    ].sort_values("Missing Values", ascending=False)
)

print("Total missing values:", tabular_df.isnull().sum().sum())

```

```

Empty DataFrame
Columns: [Missing Values, Percentage]
Index: []

Total missing values: 0

duplicate_count = tabular_df.duplicated().sum()

print("Number of duplicate rows:", duplicate_count)

Number of duplicate rows: 0

class_counts = tabular_df["diagnosis_label"].value_counts()

print(class_counts)
print("\nClass percentages:")
print(tabular_df["diagnosis_label"].value_counts(normalize=True) *
100)

diagnosis_label
benign      357
malignant   212
Name: count, dtype: int64

Class percentages:
diagnosis_label
benign      62.741652
malignant   37.258348
Name: proportion, dtype: float64


plt.figure(figsize=(7, 5))

sns.countplot(
    data=tabular_df,
    x="diagnosis_label"
)

plt.title("Distribution of Breast Cancer Diagnosis")
plt.xlabel("Diagnosis")
plt.ylabel("Number of Samples")


plt.show()

```



```
selected_features = [  
    "mean radius",  
    "mean texture",  
    "mean perimeter",  
    "mean area",  
    "mean smoothness"  
]  
  
tabular_df[selected_features].hist(  
    figsize=(15, 10),  
    bins=25  
)  
  
plt.suptitle("Distribution of Selected Medical Features")  
plt.tight_layout()  
plt.show()
```

Distribution of Selected Medical Features




```
plt.figure(figsize=(9, 6))

sns.boxplot(
    data=tabular_df,
    x="diagnosis_label",
    y="mean radius"
)

plt.title("Mean Radius by Diagnosis")
plt.xlabel("Diagnosis")
plt.ylabel("Mean Radius")

plt.show()
```



```
malignant_radius = tabular_df.loc[
    tabular_df["diagnosis_label"] == "malignant",
    "mean radius"
]

benign_radius = tabular_df.loc[
    tabular_df["diagnosis_label"] == "benign",
    "mean radius"
]

t_stat, p_value = stats.ttest_ind(
    malignant_radius,
    benign_radius,
    equal_var=False
)

print("T-statistic:", t_stat)
print("P-value:", p_value)

alpha = 0.05

if p_value < alpha:
    print(
```

```

        "\nResult: Reject the null hypothesis."
    )
    print(
        "There is a statistically significant difference "
        "in mean radius between the two diagnosis groups."
    )
else:
    print(
        "\nResult: Fail to reject the null hypothesis."
    )
    print(
        "There is insufficient evidence of a significant "
        "difference in mean radius."
    )

```

T-statistic: 22.208797758464527
P-value: 1.6844591259582815e-64

Result: Reject the null hypothesis.
There is a statistically significant difference in mean radius between the two diagnosis groups.

```

tabular_df["radius_category"] = pd.qcut(
    tabular_df["mean radius"],
    q=3,
    labels=["Low", "Medium", "High"]
)

```

```

contingency_table = pd.crosstab(
    tabular_df["radius_category"],
    tabular_df["diagnosis_label"]
)

```

```

print("Contingency Table:")
display(contingency_table)

```

```

chi2, p, dof, expected = stats.chi2_contingency(contingency_table)

```

```

print("Chi-square statistic:", chi2)
print("P-value:", p)
print("Degrees of freedom:", dof)

```

```

if p < 0.05:
    print(
        "\nReject the null hypothesis: "
        "radius category and diagnosis are associated."
    )
else:
    print(
        "\nFail to reject the null hypothesis: "

```

```

        "there is insufficient evidence of an association."
    )

contingency_table.plot(
    kind="bar",
    figsize=(9, 6)
)

plt.title("Diagnosis Distribution Across Radius Categories")
plt.xlabel("Radius Category")
plt.ylabel("Number of Samples")

plt.xticks(rotation=0)
plt.show()

```

Contingency Table:


diagnosis_label	benign	malignant
radius_category		
Low	185	6
Medium	149	39
High	23	167

Chi-square statistic: 325.4354164758156

P-value: 2.150786017083972e-71

Degrees of freedom: 2

Reject the null hypothesis: radius category and diagnosis are associated.



```
low_group = tabular_df.loc[
    tabular_df["radius_category"] == "Low",
    "mean texture"
]

medium_group = tabular_df.loc[
    tabular_df["radius_category"] == "Medium",
    "mean texture"
]

high_group = tabular_df.loc[
    tabular_df["radius_category"] == "High",
    "mean texture"
]

f_stat, p_value = stats.f_oneway(
    low_group,
    medium_group,
    high_group
)

print("F-statistic:", f_stat)
print("P-value:", p_value)
```

```

if p_value < 0.05:
    print(
        "\nReject H0: At least one group's mean texture differs
significantly."
    )
else:
    print(
        "\nFail to reject H0: No statistically significant difference
detected."
    )

```

F-statistic: 29.075477527878192
P-value: 9.551730890118103e-13

Reject H₀: At least one group's mean texture differs significantly.

```

X_tabular = tabular_df.drop(
    columns=[
        "diagnosis",
        "diagnosis_label",
        "radius_category"
    ]
)

y_tabular = tabular_df["diagnosis"]

X_train, X_test, y_train, y_test = train_test_split(
    X_tabular,
    y_tabular,
    test_size=0.20,
    stratify=y_tabular,
    random_state=RANDOM_STATE
)

print("Training shape:", X_train.shape)
print("Testing shape:", X_test.shape)

scaler = StandardScaler()

X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

X_train_tensor = X_train_scaled.astype(np.float32)
X_test_tensor = X_test_scaled.astype(np.float32)

print("X_train tensor shape:", X_train_tensor.shape)
print("X_test tensor shape:", X_test_tensor.shape)
print("Data type:", X_train_tensor.dtype)

tabular_engineered = X_tabular.copy()

```

```
tabular_engineered["radius_texture_interaction"] = (
    tabular_engineered["mean radius"]
    * tabular_engineered["mean texture"]
)
```

```
tabular_engineered["compactness_ratio"] = (
    tabular_engineered["mean compactness"]
    /
    (tabular_engineered["mean smoothness"] + 1e-8)
)
```

```
tabular_engineered.head()
```

Training shape: (455, 30)

Testing shape: (114, 30)

X_train tensor shape: (455, 30)

X_test tensor shape: (114, 30)

Data type: float32

	mean radius	mean texture	mean perimeter	mean area	mean smoothness \
0	17.99	10.38	122.80	1001.0	0.11840
1	20.57	17.77	132.90	1326.0	0.08474
2	19.69	21.25	130.00	1203.0	0.10960
3	11.42	20.38	77.58	386.1	0.14250
4	20.29	14.34	135.10	1297.0	0.10030

	mean compactness	mean concavity	mean concave points	mean symmetry \
0	0.27760	0.3001	0.14710	0.2419
1	0.07864	0.0869	0.07017	0.1812
2	0.15990	0.1974	0.12790	0.2069
3	0.28390	0.2414	0.10520	0.2597
4	0.13280	0.1980	0.10430	0.1809

	mean fractal dimension	...	worst perimeter	worst area	worst smoothness \
0	0.07871	...	184.60	2019.0	0.1622

1	0.05667	...	158.80	1956.0
0.1238				
2	0.05999	...	152.50	1709.0
0.1444				
3	0.09744	...	98.87	567.7
0.2098				
4	0.05883	...	152.20	1575.0
0.1374				

	worst compactness	worst concavity	worst concave points	worst symmetry \
0	0.6656	0.7119	0.2654	
0.4601				
1	0.1866	0.2416	0.1860	
0.2750				
2	0.4245	0.4504	0.2430	
0.3613				
3	0.8663	0.6869	0.2575	
0.6638				
4	0.2050	0.4000	0.1625	
0.2364				

	worst fractal dimension	radius_texture_interaction compactness_ratio
0	0.11890	186.7362
2.344594		
1	0.08902	365.5289
0.928015		
2	0.08758	418.4125
1.458941		
3	0.17300	232.7396
1.992281		
4	0.07678	290.9586
1.324028		

[5 rows x 32 columns]

```

selector = SelectKBest(
    score_func=f_classif,
    k=10
)

X_selected = selector.fit_transform(
    X_train_scaled,
    y_train
)

selected_feature_names = X_train.columns[
    selector.get_support()
]

```

```

print("Selected Features:")
print(selected_feature_names.tolist())

print("\nOriginal number of features:", X_train.shape[1])
print("Selected number of features:", X_selected.shape[1])

Selected Features:
['mean radius', 'mean perimeter', 'mean area', 'mean concavity', 'mean
concave points', 'worst radius', 'worst perimeter', 'worst area',
'worst concavity', 'worst concave points']

Original number of features: 30
Selected number of features: 10

```

Textual Data

```

medical_text_ds = load_dataset(
    "bhargavi909/Medical_Transcriptions"
)

print(medical_text_ds)

text_df = medical_text_ds["train"].to_pandas()

print("Shape:", text_df.shape)

text_df.head()

Warning: You are sending unauthenticated requests to the HF Hub.
Please set a HF_TOKEN to enable higher rate limits and faster
downloads.

DatasetDict({
  train: Dataset({
    features: ['description', 'medical_specialty'],
    num_rows: 3999
  })
  test: Dataset({
    features: ['description', 'medical_specialty'],
    num_rows: 1000
  })
})
Shape: (3999, 2)

              description
medical_specialty
0  Bilateral Mammogram, (abnormal) additional vi...
33

```

```

1   Bronchoscopy with brush biopsies. Persistent...
3
2   This is a middle-aged female with two month h...
22
3   MRI Brain: Thrombus in torcula of venous sinu...
33
4   Followup of moderate-to-severe sleep apnea. ...
15

print(text_df.columns.tolist())

text_df.info()

for column in text_df.columns:
    print(f"\nCOLUMN: {column}")
    print(text_df[column].head(3))

TEXT_COLUMN = "description"

['description', 'medical_specialty']
<class 'pandas.DataFrame'>
RangeIndex: 3999 entries, 0 to 3998
Data columns (total 2 columns):
#   Column                Non-Null Count  Dtype
---  -
0   description           3999 non-null   str
1   medical_specialty     3999 non-null   int64
dtypes: int64(1), str(1)
memory usage: 578.5 KB

COLUMN: description
0   Bilateral Mammogram, (abnormal) additional vi...
1   Bronchoscopy with brush biopsies. Persistent...
2   This is a middle-aged female with two month h...
Name: description, dtype: str

COLUMN: medical_specialty
0    33
1     3
2    22
Name: medical_specialty, dtype: int64

def clean_medical_text(text):
    """
    Basic demonstration cleaning function.

    This is NOT sufficient for production PHI de-identification.
    """

    if pd.isna(text):
        return ""

```

```

text = str(text)

# Remove email-like strings
text = re.sub(
    r"\b[A-Za-z0-9._%+-]+@[A-Za-z0-9.-]+\.[A-Za-z]{2,}\b",
    "[REDACTED_EMAIL]",
    text
)

# Remove phone-like strings
text = re.sub(
    r"\b(?:\+?1[-.\s]?)(?(\d{3})\)?[-.\s]?(\d{3})[-.\s]?(\d{4})\b",
    "[REDACTED_PHONE]",
    text
)

# Remove common date formats
text = re.sub(
    r"\b\d{1,2}[/-]\d{1,2}[/-]\d{2,4}\b",
    "[REDACTED_DATE]",
    text
)

# Remove excessive whitespace
text = re.sub(r"\s+", " ", text)

return text.strip().lower()

text_df["clean_text"] = text_df[TEXT_COLUMN].apply(
    clean_medical_text
)

text_df[["clean_text"]].head()

clean_text
0  bilateral mammogram, (abnormal) additional vie...
1  bronchoscopy with brush biopsies. persistent p...
2  this is a middle-aged female with two month hi...
3  mri brain: thrombus in torcula of venous sinuses.
4  followup of moderate-to-severe sleep apnea. th...

print("Missing original text:")
print(text_df[TEXT_COLUMN].isna().sum())

print("\nEmpty cleaned text:")
print((text_df["clean_text"].str.len() == 0).sum())

text_df = text_df[
    text_df["clean_text"].str.len() > 0
]

```

```

].copy()

print("New shape:", text_df.shape)

Missing original text:
0

Empty cleaned text:
6

New shape: (3993, 3)

text_df["character_count"] = text_df["clean_text"].str.len()

text_df["word_count"] = (
    text_df["clean_text"]
    .str.split()
    .str.len()
)

display(
    text_df[
        ["character_count", "word_count"]
    ].describe()
)

plt.figure(figsize=(9, 5))

sns.histplot(
    text_df["word_count"],
    bins=50
)

plt.title("Distribution of Medical Transcription Length")
plt.xlabel("Number of Words")
plt.ylabel("Frequency")

plt.show()

```

	character_count	word_count
count	3993.000000	3993.000000
mean	130.349111	18.503130
std	79.164268	11.937381
min	14.000000	1.000000
25%	68.000000	9.000000
50%	116.000000	16.000000
75%	179.000000	25.000000
max	491.000000	76.000000




Image Dataset

```
train_dataset = BreastMNIST(
    split="train",
    download=True
)

val_dataset = BreastMNIST(
    split="val",
    download=True
)

test_dataset = BreastMNIST(
    split="test",
    download=True
)

info = INFO["breastmnist"]

print("Dataset metadata:\n")

for key, value in info.items():
    print(f"{key}: {value}")

print("Class labels:")
print(info["label"])
```

Dataset metadata:

```
python_class: BreastMNIST
description: The BreastMNIST is based on a dataset of 780 breast
ultrasound images. It is categorized into 3 classes: normal, benign,
and malignant. As we use low-resolution images, we simplify the task
into binary classification by combining normal and benign as positive
and classifying them against malignant as negative. We split the
source dataset with a ratio of 7:1:2 into training, validation and
test set. The source images of 1×500×500 are resized into 1×28×28.
url: https://zenodo.org/records/10519652/files/breastmnist.npz?
download=1
MD5: 750601b1f35ba3300ea97c75c52ff8f6
url_64: https://zenodo.org/records/10519652/files/breastmnist_64.npz?
download=1
MD5_64: 742edef2a1fd1524b2efff4bd7ba9364
url_128:
https://zenodo.org/records/10519652/files/breastmnist_128.npz?
download=1
MD5_128: 363e4b3f8d712e9b5de15470a2aaadf1
url_224:
https://zenodo.org/records/10519652/files/breastmnist_224.npz?
download=1
MD5_224: b56378a6eefa9fed602bb16d192d4c8b
task: binary-class
label: {'0': 'malignant', '1': 'normal, benign'}
n_channels: 1
n_samples: {'train': 546, 'val': 78, 'test': 156}
license: CC BY 4.0
Class labels:
{'0': 'malignant', '1': 'normal, benign'}

print("Training images shape:")
print(train_dataset.imgs.shape)

print("\nTraining labels shape:")
print(train_dataset.labels.shape)

print("\nImage data type:")
print(train_dataset.imgs.dtype)

print("\nPixel range:")
print(
    train_dataset.imgs.min(),
    "to",
    train_dataset.imgs.max()
)

Training images shape:
(546, 28, 28)
```

Training labels shape:
(546, 1)

Image data type:
uint8

Pixel range:
0 to 255

```
fig, axes = plt.subplots(
    2,
    5,
    figsize=(12, 5)
)

for i, ax in enumerate(axes.flat):
    image, label = train_dataset[i]

    ax.imshow(
        image,
        cmap="gray"
    )


    ax.set_title(
        f"Label: {label[0]}"
    )

    ax.axis("off")

plt.suptitle(
    "Sample Images from BreastMNIST",
    fontsize=14
)

plt.tight_layout()
plt.show()
```

Sample Images from BreastMNIST



```
y_train = train_dataset.labels.flatten()
y_val = val_dataset.labels.flatten()
y_test = test_dataset.labels.flatten()
```

```
print("Training labels:", y_train.shape)
print("Validation labels:", y_val.shape)
print("Test labels:", y_test.shape)
```

```
Training labels: (546,)
Validation labels: (78,)
Test labels: (156,)
```

```
class_counts = pd.Series(
    y_train
).value_counts().sort_index()
```

```
class_percentages = (
    class_counts / class_counts.sum()
) * 100
```

```
class_summary = pd.DataFrame({
    "Count": class_counts,
    "Percentage": class_percentages
})
```

```
display(class_summary)
```

	Count	Percentage
0	147	26.923077
1	399	73.076923

```
plt.figure(figsize=(7, 5))
```

```
plt.bar(
```

```

    class_counts.index.astype(str),
    class_counts.values
)

plt.title(
    "Class Distribution in BreastMNIST Training Data"
)

plt.xlabel("Class")
plt.ylabel("Number of Images")

plt.show()

```



```

X_train_raw = train_dataset.imgs

invalid_pixels = np.sum(
    (X_train_raw < 0)
    | (X_train_raw > 255)
)

print("Invalid pixel count:", invalid_pixels)

```

```

Invalid pixel count: 0
plt.figure(figsize=(9, 5))


plt.hist(
    X_train_raw.flatten(),
    bins=50
)

plt.title(
    "Distribution of Pixel Intensities"
)

plt.xlabel("Pixel Intensity")
plt.ylabel("Frequency")

plt.show()

```



```

class_0_images = X_train_raw[
    y_train == 0
]


class_1_images = X_train_raw[
    y_train == 1
]

mean_class_0 = np.mean(


```

```
    class_0_images,  
    axis=0  
)  
  
mean_class_1 = np.mean(  
    class_1_images,  
    axis=0  
)  
  
fig, axes = plt.subplots(  
    1,  
    2,  
    figsize=(10, 5)  
)  
  
axes[0].imshow(  
    mean_class_0.squeeze(),  
    cmap="gray"  
)  
  
axes[0].set_title(  
    "Mean Image – Class 0"  
)  
  
axes[0].axis("off")  
  
axes[1].imshow(  
    mean_class_1.squeeze(),  
    cmap="gray"  
)  
  
axes[1].set_title(  
    "Mean Image – Class 1"  
)  
  
axes[1].axis("off")  
  
plt.tight_layout()  
plt.show()
```

Mean Image — Class 0



Mean Image — Class 1



```
X_train = train_dataset.imgs.astype(np.float32)
X_val = val_dataset.imgs.astype(np.float32)
X_test = test_dataset.imgs.astype(np.float32)
```

```
X_train /= 255.0
X_val /= 255.0
X_test /= 255.0
```

```
print("Training pixel range:")
print(
    X_train.min(),
    "to",
    X_train.max()
)
```

```
Training pixel range:
0.0 to 1.0
```

```
sample_image = X_train[0].squeeze()
```

```
median_denoised = median_filter(
    sample_image,
    size=3
)
```

```
fig, axes = plt.subplots(
    1,
    2,
    figsize=(9, 4)
)
```

```

axes[0].imshow(
    sample_image,
    cmap="gray"
)

axes[0].set_title("Original")
axes[0].axis("off")


axes[1].imshow(
    median_denoised,
    cmap="gray"
)

axes[1].set_title("Median Filtered")
axes[1].axis("off")


plt.tight_layout()
plt.show()

```

Original



Median Filtered



```

def extract_image_features(image):
    image = image.squeeze()

    # Calculate horizontal and vertical edges
    edge_x = sobel(
        image,
        axis=0
    )

    edge_y = sobel(

```

```

        image,
        axis=1
    )

    edge_magnitude = np.sqrt(
        edge_x**2
        +
        edge_y**2
    )

    return {
        "mean_intensity": np.mean(image),
        "std_intensity": np.std(image),
        "min_intensity": np.min(image),
        "max_intensity": np.max(image),
        "median_intensity": np.median(image),
        "q25_intensity": np.percentile(image, 25),
        "q75_intensity": np.percentile(image, 75),
        "contrast": np.max(image) - np.min(image),
        "energy": np.sum(image**2),
        "mean_edge_strength": np.mean(edge_magnitude),
        "max_edge_strength": np.max(edge_magnitude)
    }

train_features = pd.DataFrame([
    extract_image_features(image)
    for image in X_train
])

display(train_features.head())

display(
    train_features.describe().T
)

```

	mean_intensity	std_intensity	min_intensity	max_intensity	\
0	0.398204	0.152389	0.082353	0.792157	
1	0.233368	0.119642	0.007843	0.623529	
2	0.372854	0.123865	0.125490	0.811765	
3	0.425920	0.182939	0.082353	0.874510	
4	0.383523	0.165781	0.003922	0.823529	

	median_intensity	q25_intensity	q75_intensity	contrast	energy	\
0	0.400000	0.278431	0.521569	0.709804	142.522659	
1	0.229412	0.141176	0.301961	0.615686	53.919586	
2	0.360784	0.274510	0.458824	0.686275	121.020432	

3	0.474510	0.266667	0.557843	0.792157
168.461792				
4	0.364706	0.266667	0.470588	0.819608
136.865707				

	mean_edge_strength	max_edge_strength
0	0.455046	1.431346
1	0.395428	2.086407
2	0.384229	1.608981
3	0.395050	2.094837
4	0.510351	2.929010

	count	mean	std	min	25%
\					
mean_intensity	546.0	0.327581	0.084759	0.118157	0.265160
std_intensity	546.0	0.182824	0.041622	0.085227	0.152864
min_intensity	546.0	0.038691	0.040752	0.000000	0.003922
max_intensity	546.0	0.763348	0.108103	0.466667	0.694118
median_intensity	546.0	0.315528	0.111376	0.054902	0.235294
q25_intensity	546.0	0.165767	0.094596	0.003922	0.094118
q75_intensity	546.0	0.472176	0.105630	0.184314	0.392402
contrast	546.0	0.724657	0.105476	0.431373	0.650980
energy	546.0	117.312973	48.505951	21.448368	79.783745
mean_edge_strength	546.0	0.399996	0.073235	0.239485	0.352358
max_edge_strength	546.0	1.665943	0.465895	0.865628	1.356775

	50%	75%	max
mean_intensity	0.323382	0.389760	0.665966
std_intensity	0.186288	0.212189	0.288046
min_intensity	0.025490	0.058824	0.301961
max_intensity	0.772549	0.838235	1.000000
median_intensity	0.305882	0.398039	0.678431
q25_intensity	0.152941	0.227451	0.592157
q75_intensity	0.482353	0.545098	0.757843
contrast	0.729412	0.796078	0.996078
energy	112.900681	151.606674	359.839233
mean_edge_strength	0.394311	0.446844	0.712052
max_edge_strength	1.542922	1.842567	3.834147

```

train_features["intensity_range"] = (
    train_features["max_intensity"]
    -
    train_features["min_intensity"]
)

train_features["interquartile_range"] = (
    train_features["q75_intensity"]
    -
    train_features["q25_intensity"]
)

train_features["edge_to_intensity_ratio"] = (
    train_features["mean_edge_strength"]
    /
    (train_features["mean_intensity"] + 1e-8)
)

```

```
display(train_features.head())
```

	mean_intensity	std_intensity	min_intensity	max_intensity	\
0	0.398204	0.152389	0.082353	0.792157	
1	0.233368	0.119642	0.007843	0.623529	
2	0.372854	0.123865	0.125490	0.811765	
3	0.425920	0.182939	0.082353	0.874510	
4	0.383523	0.165781	0.003922	0.823529	

	median_intensity	q25_intensity	q75_intensity	contrast	\
energy					
0	0.400000	0.278431	0.521569	0.709804	
142.522659					
1	0.229412	0.141176	0.301961	0.615686	
53.919586					
2	0.360784	0.274510	0.458824	0.686275	
121.020432					
3	0.474510	0.266667	0.557843	0.792157	
168.461792					
4	0.364706	0.266667	0.470588	0.819608	
136.865707					

	mean_edge_strength	max_edge_strength	intensity_range	\
0	0.455046	1.431346	0.709804	
1	0.395428	2.086407	0.615686	
2	0.384229	1.608981	0.686275	
3	0.395050	2.094837	0.792157	
4	0.510351	2.929010	0.819608	

	interquartile_range	edge_to_intensity_ratio
0	0.243137	1.142745
1	0.160784	1.694436

2	0.184314	1.030509
3	0.291176	0.927520
4	0.203922	1.330689

```
def augment_image(image):  
    # Remove channel dimension temporarily  
    image = image.squeeze()  
  
    # Small random rotation  
    angle = np.random.uniform(-5, 5)  
  
    augmented = rotate(  
        image,  
        angle=angle,  
        reshape=False,  
        mode="nearest"  
    )  
  
    # Small random horizontal/vertical shift  
    shift_amount = np.random.uniform(-1, 1, size=2)  
  
    augmented = shift(  
        augmented,  
        shift=shift_amount,  
        mode="nearest"  
    )  
  
    # Ensure valid normalized range  
    augmented = np.clip(  
        augmented,  
        0,  
        1  
    )  
  
    return augmented.astype(np.float32)  
  
original = X_train[4].squeeze()  
  
augmented = augment_image(  
    X_train[4]  
)  
  
fig, axes = plt.subplots(  
    1,  
    2,  
    figsize=(9, 4)  
)  
  
axes[0].imshow(  

```

```

        original,
        cmap="gray"
    )

    axes[0].set_title("Original Image")
    axes[0].axis("off")


    axes[1].imshow(
        augmented,
        cmap="gray"
    )

    axes[1].set_title("Augmented Image")
    axes[1].axis("off")


    plt.tight_layout()
    plt.show()

```

Original Image



Augmented Image



Signal Data

```

ecg = electrocardiogram()

sampling_rate = 360

print("ECG shape:", ecg.shape)
print("Number of samples:", len(ecg))
print("Duration in seconds:", len(ecg) / sampling_rate)

```

```
Downloading file 'ecg.dat' from  
'https://raw.githubusercontent.com/scipy/dataset-ecg/main/ecg.dat' to  
'/home/lakshya/.cache/scipy-data'.
```

```
ECG shape: (108000,)
```

```
Number of samples: 108000
```

```
Duration in seconds: 300.0
```

```
time = np.arange(len(ecg)) / sampling_rate
```

```
plt.figure(figsize=(14, 5))
```


```
plt.plot(  
    time[:10 * sampling_rate],  
    ecg[:10 * sampling_rate]  
)
```

```
plt.title("ECG Signal – First 10 Seconds")
```

```
plt.xlabel("Time (seconds)")
```

```
plt.ylabel("Amplitude (mV)")
```

```
plt.show()
```



```
def bandpass_filter(  
    signal,  
    lowcut,  
    highcut,  
    fs,  
    order=4  
):  
    nyquist = 0.5 * fs  
    low = lowcut / nyquist
```

```

    high = highcut / nyquist

    b, a = butter(
        order,
        [low, high],
        btype="band"
    )

    filtered_signal = filtfilt(
        b,
        a,
        signal
    )

    return filtered_signal

ecg_filtered = bandpass_filter(
    ecg,
    lowcut=0.5,
    highcut=40,
    fs=sampling_rate
)

start = 9 * sampling_rate
end = int(10.2 * sampling_rate)

plt.figure(figsize=(14, 5))


plt.plot(
    time[start:end],
    ecg[start:end],
    label="Original ECG",
    alpha=0.7
)

plt.plot(
    time[start:end],
    ecg_filtered[start:end],
    label="Filtered ECG"
)

plt.title("Original vs Filtered ECG")
plt.xlabel("Time (seconds)")
plt.ylabel("Amplitude (mV)")

plt.legend()
plt.show()

```



```
ecg_mean = ecg_filtered.mean()
ecg_std = ecg_filtered.std()

ecg_normalized = (
    ecg_filtered - ecg_mean
) / ecg_std

print("Mean after normalization:", ecg_normalized.mean())
print("Standard deviation:", ecg_normalized.std())
```

```
Mean after normalization: -1.5954316057576323e-18
Standard deviation: 1.0
```

```
WINDOW_SECONDS = 5
```

```
WINDOW_SIZE = (
    WINDOW_SECONDS * sampling_rate
)
```

```
NUM_WINDOWS = len(ecg_normalized) // WINDOW_SIZE
```

```
signal_windows = ecg_normalized[
    :NUM_WINDOWS * WINDOW_SIZE
].reshape(
    NUM_WINDOWS,
    WINDOW_SIZE
)
```

```
signal_windows = signal_windows.astype(np.float32)
```

```
print("Signal window tensor shape:")
print(signal_windows.shape)
```

```
Signal window tensor shape:
(60, 1800)
```

```
def extract_signal_features(window):
    return {
        "mean": np.mean(window),
        "std": np.std(window),
        "min": np.min(window),
        "max": np.max(window),
        "range": np.ptp(window),
        "energy": np.sum(window ** 2),
        "mean_absolute_value": np.mean(np.abs(window))
    }

signal_features = pd.DataFrame(
    [
        extract_signal_features(window)
        for window in signal_windows
    ]
)

display(signal_features.head())

display(signal_features.describe())
```

	mean	std	min	max	range	energy \
0	0.023410	0.864440	-1.323456	5.092521	6.415977	1346.046997
1	-0.006466	0.673005	-1.056962	3.908225	4.965186	815.358948
2	-0.003934	0.652378	-1.048052	3.587795	4.635848	766.102966
3	0.008633	1.126776	-1.907000	5.532056	7.439055	2285.458252
4	0.010813	0.980489	-2.013973	5.128770	7.142743	1730.656372

	mean_absolute_value
0	0.519758
1	0.412980
2	0.372546
3	0.682996
4	0.616996

	mean	std	min	max	range
energy \					
count	6.000000e+01	60.000000	60.000000	60.000000	60.000000
mean	4.501392e-10	0.987238	-1.974595	5.097237	7.071832
std	2.538208e-02	0.158577	0.598912	0.790185	1.014425
min	-7.090084e-02	0.469738	-4.365462	2.557551	4.231677
25%	-1.347381e-02	0.922772	-2.204772	4.475148	6.411532
50%	-1.110820e-04	0.993267	-2.015231	5.154497	7.189032

```
1777.013916
75%      1.081733e-02    1.079110  -1.703412    5.726050    7.811046
2096.536865
max      6.749131e-02    1.396435  -0.726895    6.526565    9.017020
3510.268799
```

```
      mean_absolute_value
count          60.000000
mean           0.606415
std            0.119245
min            0.274387
25%            0.550354
50%            0.596661
75%            0.676580
max            0.965590
```